

Average Reward Adjusted Deep Reinforcement Learning: Near-Blackwell-Optimal Policies applied to the Order Release Problem

Manuel Schneckenreither¹, Stefan Haeussler¹, and Juanjo Peiró²

¹ Department of Information Systems, Production and Logistics Management,
University of Innsbruck, Austria

`manuel.schneckenreither@uibk.ac.at, stefan.haeussler@uibk.ac.at`

² Departament d'Estadística i Investigació Operativa, Universitat de València, Spain
`juanjo.peiro@uv.es`

Abstract. A major task in hierarchically structured manufacturing planning and control is to determine when to release orders to the shop floor. One key parameter is the lead time which is the planned time that elapses between the release of an order and its completion. Lead times are usually determined based on the flow times, i.e. the measured duration orders take to traverse the production system. Traditional order release models assume static lead times, although it has been shown that they should be set dynamically to be able to react to the dynamic operational characteristics of production system. Therefore, we present a reinforcement learning order release model which sets lead times dynamically. The applied reinforcement learning agent is designed especially for processes with periodic feedback and high variability. This paper aims at providing a proof-of-concept for the proposed algorithm. We compare our method on a three-stage flow-shop simulation model not only to static order release techniques, but also to the commonly used deep Q-Learning agent. The results show that both reinforcement learning agents have room for improvement for queuing systems with low variability, that our proposed method usually outperforms deep Q-Learning and finally for systems with high utilisation that the proposed reinforcement learning algorithm performs very well, outperforming all other approaches.

Keywords: operations research, production planning, order release, machine learning, reinforcement learning

1 Introduction

An important goal in Manufacturing Planning and Control (MPC) systems is to achieve short and predictable flow times, especially where high flexibility in meeting customer demand is required. Besides achieving short flow times, one should also maintain high output and due-date performance while keeping the work-in-process level low. One approach to address this problem is to collect all incoming orders in an order-pool and periodically decide which orders to release

to the shop floor. Once orders are released, costs start to accumulate as planned orders materialise as actual jobs in the production system. The main challenge is to find a good compromise of balancing the shop floor and timely completion of jobs. Although the performance of such systems can be measured manifold the most overarching objective is to maximise profits by adequately assigning holding and lateness costs.

One of the key modeling parameters for order release mechanisms is the *lead time*, which refers to the planned time that elapses between the release of an order and its arrival in the finished goods inventory. This planning parameter is often based on the observable time an order needs to traverse the production system, which in contrast is denoted as the *flow time*. Flow times consist of processing, setup, control, transport, and waiting times, whereas the latter is the governing factor (Zäpfel, 1982, p.223). Waiting times are a result from queuing (e.g. jobs queue before and after processing), depend heavily on the amount of jobs in the system (WIP) and thus are relatively difficult to estimate, which makes the setting of favorable lead times so difficult (e.g., Tatsiopoulos and Kingsman 1983; Wiendahl 1995). Most state-of-the-art order release mechanisms use static or fixed lead times to address the order release problem, and thus neglect the nonlinear relationship between resource utilisation and flow times, which is well known from practice and queuing theory (Pahl, Voss, and Woodruff 2007).

One way to address this nonlinear interaction effects is to set lead times dynamically. Intuitively the order release problem is solved by perfectly matching the lead times to the flow times, but the corresponding optimisation problem faces “sampling issues” meaning that the flow times depend on the lead times. Thus, a completely reactive approach can lead to an erratic order release behaviour, the Lead Time Syndrome (see also Knollmann and Windt 2013; Mather and Plossl 1978; Selcuk, Fransoo, and De Kok 2006). We hypothesise that dynamic lead time management approaches require a predictive design, i.e. the method has to incorporate a prediction of the future system state, to be able to compete with state-of-the-art order release methods from literature. This is necessary in order to detect issues, like the Lead Time Syndrome, in advance and react accordingly (Schneckenreither, Haeussler, and Gerhold, 2020). However, there only exist three papers that propose a predictive lead time management algorithm in literature.

One method for which the future system state is by design considered a main driver of decision making is reinforcement learning (e.g. Sutton, Barto et al., 1998), one of the main three subcategories of machine learning. Reinforcement learning (RL) is an optimisation technique that stems from dynamic programming, therefore also seeks to compute the best stationary policy for a given problem. The idea is to evaluate the current system state according to the direct feedback (reward), as well as the previously learned state value of the consecutive state. This forms a cyclic constraint problem with as many constraints and variables as there are system states. The solution is found through iteration and due to the transitivity property of the state value update. The advantage of RL over dynamic programming is that (i) the problem space is explored by an

agent and thus only expectantly interesting parts of the problem space need to be assessed and (ii) the knowledge (acquisition) of transition probabilities becomes unnecessary as the states are evaluated by consecutively observed states solely.

There exist two main types of deep RL algorithms in literature. The commonly applied discounted RL method (e.g. Mnih et al., 2016, 2015; Sutton, Barto et al., 1998), in which the future system state is discounted by a factor $0 < \gamma < 1$ to prevent an infinite growth of the state evaluations, and average reward RL (Mahadevan, 1996b; Mahadevan et al., 1997; Schwartz, 1993). The advantage of the later is that the underlying optimisation technique is theoretically able to find better policies by splitting up the evaluation of each state into infinitely many components. First, the average reward per step, second, a bias value that specifies expected reward that is additionally accumulated when traversing from the current state to the optimal path, and third, an infinite number of values of which each includes the occurrence of the rewards and in which magnitude along the path (Howard, 1964; Mahadevan, 1996b; Puterman, 1994). To expectantly accumulate the most reward one has to choose the action that leads to the state with highest value of the components following the order above, on ties selecting according to the next component. That is, the highest evaluation of the average reward, the bias value, the third component, the fourth component, etc. The formulations for computing these values form nested constraints, so for instance the current estimate of the average reward is used to compute the bias value. Therefore, a slight approximation error may cause harmful damage in subsequent constraints, which makes average reward reinforcement computationally very intensive. Worsening the issue, Puterman (1994, p.343ff) shows that the constraint problem has to be specified to solve one more component as originally intended to gain correct solution.

In contrast the standard discounted framework estimates only one single scalar value consisting of the addition of these components, where however, the average reward is scaled by $1/(1 - \gamma)$. But in order to collect as much reward as possible the discount factor γ has to be very close to one, for instance 0.99. Such a parameterisation leads to the fact that the subterm including the average reward dominates the other two. Further, in commonly applied unichain (and thus ergodic) MDPs the average reward per step is equal over all states. Thus, independently assessing it is not only computationally unwisely, but also problematic as they are iteratively approximated. For this reasons most applications of discounted RL in literature evaluate their algorithms on processes that have an average reward of about zero (Schneckenreither, 2020). Schneckenreither (2020) shows that tabular based discounted RL agents are unable to find even basic policies when confronted with processes that have a small average reward per step. Unfortunately, when optimising for the highest profit (or the lowest cost) in the area of operations management, the most intuitive and clear way of reward feedback to the agent is presenting it the profit that is accumulated every period.

Therefore, we develop an algorithm that uses only one scalar value for the estimation of the average reward over all states. Further it estimates for every state-action pair the bias value and the sum of all other state value components, usually referred to as error term. The later values are adjusted by the average reward. Thus, as opposed to the commonly applied discounted RL algorithm we use an average reward adjusted RL algorithm to dynamically set lead time for order release. In contrast to the aforementioned works on average reward RL our algorithm incorporates the optimisation of not only the average reward over time, but also the bias values, which is an important adaption in highly stochastic systems. Only the work by Mahadevan (1996a) integrates this second-level refinement optimisation. However, their average reward reinforcement learning algorithm requires the selection of a reference state to prevent an infinite increase of state values. This results from the lack of feedback in the iterative process of optimising different sets of nested constraints. Furthermore, they assess the average reward independently for each state.

To achieve a scalable lead time setting algorithm, we propose an extension that reduces the action space tremendously. In fact, it transforms the problem size function measured on the number of product types from one with an exponential increase to a linear one. The idea is that multiple action decision agents are introduced that operate on the same state, but each one is responsible for the lead time setting of only one product type.

Finally, we provide a proof-of-concept experimental section by testing the proposed algorithm on a three stage flow shop with six product types, which provides ample evidence of the viability of the approach. The highly optimised value-iteration algorithm with improved n-step learning, small action-based replay memories and concurrent workers infers order release policies which, especially in case of high utilisation, outperform not only the static order release mechanisms, but also the commonly applied deep Q-Learning algorithm (cf. Mnih et al., 2015). In summary our contributions are as follows:

- We propose a novel *average reward adjusted deep reinforcement learning algorithm* ARAL. This is the first average reward adjusted reward learning algorithm that operates using an artificial neural network as function approximation. It includes several known extensions like n-step learning (e.g. Mnih et al., 2016) or experience replay memory (e.g. Mnih et al., 2015), which we additionally advanced to prevent catastrophic forgetting by action-based selection of experiences.
- We establish an action reduction method, named *independent agents*, as a subset of the action space is treated as if there were multiple independent agents, although they share the same state. It reduces the growth of the number of possible actions from an exponential growth to a linear one. This is especially useful for product type based lead time management, as each agent is responsible for the lead time of one product type.
- Using a three-stage flow shop model we show in an *experimental evaluation* that the proposed algorithm outperforms standard Q-Learning and well established order release methods, especially for systems with high variability.

Structure. The rest of the paper is structured as follows. The next section introduces related literature. Section 3 establishes the average reward adjusted deep reinforcement learning algorithm and the extension to independent agents. Section 4 describes the simulation model we use to evaluate the approach. In Section 5 we briefly discuss lead time setting in this paper, while the results of the computation experiments are presented in Section 6. Finally Section 7 concludes the paper.

2 Related Literature

This section is split into several parts. First we discuss lead time management literature, then RL applications to lead time management, third we briefly review the relevant notation to RL and finally we present the relevant literature of reinforcement learning in the rest of the section. The later includes a discussion of discounted reinforcement learning, average reward RL, the Laurent series expansion of the state values, and finally average reward adjusted RL.

2.1 Related Lead Time Management Literature

Setting lead times dynamically is a rather difficult task. Therefore, most state-of-the-art order release mechanisms use static or fixed lead times to address the order release problem. However, lead times should be set dynamically due to the non-linear relationship between utilisation and flow times (Hoyt, 1978; Pahl, Voss, and Woodruff, 2007), but this may also substantially degrade the system performance. For instance, a solely reactive approach can result in the lead time syndrome. The *lead time syndrome* describes a vicious cycle where increasing flow times perpetually inflate the lead times which leads to worse performance (see e.g., Knollmann and Windt 2013; Mather and Plossl 1978; Selcuk, Fransoo, and De Kok 2006). Schneckenreither, Haeussler, and Gerhold (2020) have established following categorisation of dynamic lead time management approaches:

- **Reactive lead time management** sets lead times by reacting on earlier flow times (e.g., Enns and Suwanruji 2004; Selcuk, Fransoo, and De Kok 2006), i.e. data from the *past*.
- **Proactive lead time management** may incorporate *past* data and the *current* system state to set lead times (e.g., Bertrand 1983; Chung and Huang 2002).
- **Predictive lead time management** may incorporate *past* data and the *current* system state, and uses an anticipated *future* system state to prevent arising issues of future periods (e.g. Paternina-Arboleda and Das 2001; Schneckenreither 2019; Schneckenreither, Haeussler, and Gerhold 2020).

In the sequel we review the lead time management (LTM) literature for due date assignment, order release and scheduling, with the focus on order release literature.

Reactive LTM. Early due date assignment approaches (e.g. Conway and Maxwell, 1967) set the due dates according to the processing time (Total Work Content rule, TWK rule) or the number of operations. For order release Enns and Suwanruji (2004) investigate the use of a reactive approach with exponential smoothing of flow times to set lead times for two-stage distribution requirements planning. Although it is considered an exploratory research paper the main conclusions are that dynamically setting planned lead times using exponential smoothing works well, but that the reactions are delayed. The latter is due to the nature of the flow time measurement only being possible for finished orders. Another reactive approach uses exponential smoothing to forecast flow times to set lead times for a capacitated multi-stage make-to-stock system (Selcuk, Fransoo, and De Kok, 2006). Their approach performs better than fixed lead time techniques for most experiments. However, they also exhibit the LTS.

Proactive LTM. The TWK due date assignment rules have been extended by several researchers to proactive methods (Eilon and Chowdhury, 1976; Enns, 1995; Kanet, 1986). As flow times depend on the workload and machine capacity in the shop (Bertrand, 1983) workload dependent due date assignment rules have been developed. The basic idea is to estimate the flow time of an order on its arrival by measuring the number of orders queued on the corresponding routing, which can be analytically (e.g. Chung and Huang, 2002), by the use of a simulation model (e.g. Vig and Dooley, 1991), statistical methods (e.g. Tai, Pearn, and Lee, 2012) or using hybrid approaches (e.g. Kaplan and Unal, 1993). Furthermore several proactive due date assignment methods that incorporate neural networks exist (Chang, Wang, and Ting, 2008; Hsu and Sha, 2004; Patil, 2008; Philipoom, Rees, and Wiegmann, 1994; Philipoom, Wiegmann, and Rees, 1997). In the case of order release the analytical papers by Selcuk et al. (2009) and Selcuk (2013) use the WIP level to quote lead times. In the former they show that the less frequently updating the lead time decreases the nervousness and improves the performance, while in the latter that dynamically updating lead times performs much better than static lead times. Modified infinite loading extends the static Backward Infinite Loading (BIL) procedure to integrate shopfloor congestion by subtracting number of queued jobs on the corresponding routing from the due date (Ragatz and Mabert, 1988; Wisner, 1995). For non-due-date oriented dispatching rules, their approach has a positive effect on total cost performance, while for due-date oriented dispatching this holds only if the due dates are loose. For all other cases the release algorithm has no negative effect on performance. Variants of this rule-based method exist (e.g. Kim and Bobrowski, 1995a; Park and Bobrowski, 1989). A machine learning approach is presented by Lee, Piramuthu, and Tsai (1997). They combine a genetic algorithm with a decision tree to generate sequencing rules in the order pool for release. In the tested job shops the CONWIP rule determines the release of the orders and a decision tree to find the order release sequence. Their method leads to significant improvements over conventional single release and dispatching approaches.

Predictive LTM. Due to the complex problem structure we hypothesise that dynamic lead time management approaches should aim for a predictive design in order to be able to compete with state-of-the-art order release methods from literature. However, there only exist four papers that propose a predictive lead time management method.

For scheduling Wang and Usher (2007) present a predictive lead time management approach using RL. Their algorithm uses multiple agents to forecast the processing time of each job for routing decisions. They show on a 10 machine job shop that the RL algorithm performs equally or better than tested heuristics. In case of order release, Paternina-Arboleda and Das (2001) introduce a predictive order release model using RL in a single product, four-station serial flow line. They compare its performance (WIP costs) with conventional order release policies (e.g., Kanban and CONWIP). The algorithm decides on whether or not to release an order after each system change, the completion of an operation of any order at any stage or a new order arrival, and assumes that any unsatisfied demand is lost. Thus, they use a continuous order release method although in practice order release decisions often need to be made on a periodical basis, e.g. daily (see Enns and Suwanruji 2004; Gelders and Van Wassenhove 1982). They outperform existing control policies with their tabular based RL agent. Then, Schneckenreither (2019) use several different RL algorithms to make periodic order release decisions for a flow shop production system. The algorithm directly sets lead times for each product type, which are then used to release the orders in the order pool. They show that their approach outperforms static order release mechanisms by yielding lower costs, lateness and standard deviation of lateness, but conclude that research using average reward reinforcement learning methods harbor optimisation potential for the order release problem. And finally, Schneckenreither, Haeussler, and Gerhold (2020) present a flow time estimation procedure to set lead times dynamically using an artificial neural network, which is used to forecast flow times. By implementing a rolling horizon order release simulation of the proceeding periods they lift their approach to a predictive lead time management approach, which is able to detect backorders of future periods and reacts by releasing orders earlier. Nonetheless, their method is unable to foresee the triggering of the lead time syndrome and therefore they introduce an upper lead time bound to prevent the negative effects of the syndrome. Their model outperforms static and reactive lead time management approaches, especially for scenarios with high utilisation.

2.2 Related Reinforcement Learning Applications

Over the past decades RL has been applied to various problems, for which astonishing results have been reported. E.g. only recently Mnih et al. (2015) have presented a novel value-iteration RL agent which exceeds human-level abilities in playing many classic Atari games (Bellemare, Veness, and Bowling, 2012). Further, Mnih et al. (2016) present improved results with asynchronous actor-critic RL. Also games like Go (Silver et al., 2016) and Chess (Silver et al., 2017) have been mastered with superhuman performance by *tabula rasa* RL agents.

Furthermore, the method has also been applied in the setting of manufacturing system, e.g. to improve the ramp-up process (Doltsinis, Ferreira, and Lohse, 2012), in locally selecting appropriate dispatching rules (Wang and Usher, 2005; Zhang and Dietterich, 1995) or scheduling (Waschneck et al., 2018; Zhang and Dietterich, 1995). However, all these applications use discounted RL and are either designed to investigate a rather simple MDP, e.g. by selecting heuristics instead of optimising the underlying problem itself, or by mapping the actual objective in a reward function that is approximately zero on average over time. This is due to the fact that state value is largely composed of a term defined by the current policies’ average reward value (Blackwell, 1962; Miller and Veinott, 1969) which would otherwise dilute the state values and thus decrease the solution quality, as can for instance be observed in Schneckenreither and Haeussler (2019) and Gijsbrechts et al. (2018).

Therefore, most applications that incorporate and directly reflect costs or profit in the reward function use average reward RL. Aydin and Öztemel (2000) use it in the setting of scheduling, while in a series of papers Mahadevan et al. investigated several problem domains starting with simple MDPs (Mahadevan, 1996b,c). After these foundational works for average reward RL they introduced a continuous time average reward RL algorithm named SMART (Mahadevan et al., 1997). Applications of SMART reach from the optimisation of queuing systems (Mahadevan, 1996a,d), maintenance of an inventory system (Das et al., 1999) to optimising transfer line in terms of maximizing demand, while keeping inventory levels of unfinished product as low as possible (Mahadevan and Theocharous, 1998). However, in practise usually decision have to be made on a daily basis (Enns and Suwanruji, 2004; Gelders and Van Wassenhove, 1982). Therefore, we refrain from this adaption and concentrate on standard MDPs only. Furthermore, often continuous-time semi-MDP problems can be converted through uniformisation into equivalent discrete time instances (see Bertsekas et al., 1995; Puterman, 1994).

2.3 Reinforcement Learning Preliminaries

The rest of this section briefly reintroduces the important concepts of average reward adjusted RL, elaborates on optimality criteria and provides insights of the underlying algorithm. For a more extensive introduction to average reward RL we refer to Schneckenreither (2020).

Like Miller and Veinott (1969) we are considering problems that are observed in a sequence of points in time labeled $1, 2, \dots$ and can be modelled using a finite set of states \mathcal{S} , labelled $1, 2, \dots, |\mathcal{S}|$, where the size $|\mathcal{S}|$ is the number of elements in \mathcal{S} . At each point t in time the system is in a state $s_t \in \mathcal{S}$. Further, by choosing an action a_t of a finite set of possible actions A_s the system returns a reward $r_t = r(s_t, a_t)$ and transitions to another state $s_{t+1} \in \mathcal{S}$ at time $t + 1$ with conditional probability $p(s_{t+1}, r_t \mid s_t, a_t)$. That is we assume that reaching state s_{t+1} from state s_t with reward r_t depends solely on the previous state s_t and chosen action a_t . In other words, we expect the system to possess the Markov property (Sutton, Barto et al., 1998, p.63). RL processes that possess the Markov

property are referred to as Markov decision processes (MDPs) (Sutton, Barto et al., 1998, p.66).

Thus, the action space is defined as $F = \times_{s=1}^{|S|} A_s$, where A_s is a finite set of possible actions. A *policy* is a sequence $\pi = (f_1, f_2, \dots)$ of elements $f_t \in F$. Using the policy π means that if the system is in state s at time t the action $f_t(s)$, i.e. the s -th component of f_t , is chosen. A stationary policy $\pi = (f, f, \dots)$ does not depend on time. In the sequel we are concerned with stationary policies only.

2.4 Discounted Reinforcement Learning

In the widely applied discounted framework the value of a state $V_\gamma^{\pi_\gamma}(s)$ is defined as the expected discounted sum of rewards under the stationary policy π_γ when starting in state s . Note that the policy π_γ depends on the selected discount factor. That is

$$V_\gamma^{\pi_\gamma}(s) = \lim_{N \rightarrow \infty} \mathbb{E} \left[\sum_{t=0}^{N-1} \gamma^t R_t^{\pi_\gamma}(s) \right],$$

where $0 \leq \gamma < 1$ is the discount factor and $R_t^\pi(s) = \mathbb{E}_\pi[r(s_t, a_t) \mid s_t = s, a_t = a]$ the reward received at time t upon starting in state s by following policy π (Mahadevan, 1996b). The aim in the discounted framework is to find an optimal policy π_γ^* , which when followed, maximises the state value for all states s as compared to any other policy π_γ : $V_\gamma^{\pi_\gamma^*} - V_\gamma^{\pi_\gamma} \geq 0$. This criteria is usually referred to as γ -optimality as the discount factor γ is fixed. However, this also means that the actual value set for γ determines the best achievable policy. For instance, as can be seen in Figure 1 setting $\gamma < 0.8027$ the printer-loop is preferred over the mail-loop, although the mail-loop accumulates more reward over time, i.e. selecting the mail-loop is a sub-optimal choice. Observe that the average reward received per step equals 1 for the printer-loop and 2 for the mail-loop. Thus, the (Blackwell-)optimal policy is to choose the mail-loop. However, if $\gamma < 3^{-\frac{1}{5}} \approx 0.8027$ an agent using discounted RL chooses the printer loop as doing the printer loop has a higher evaluation.

Therefore, in the seek of a more general optimality criteria for determining the best stationary policy π for a given problem, we introduce n -discount-optimality.

Definition 1 (n -Discount-Optimality). Due to Veinott (1969) for a MDP a policy π^* is n -discount-optimal for $n = -1, 0, 1, \dots$ for all states $s \in \mathcal{S}$ with discount factor γ if and only if

$$\lim_{\gamma \rightarrow 1} (1 - \gamma)^{-n} (V_\gamma^{\pi^*}(s) - V_\gamma^\pi(s)) \geq 0.$$

Only if a policy is n -discount-optimal for all $n < m$ it can be m -discount optimal (Puterman, 1994; Veinott, 1969). If a policy is -1 -discount-optimal it is called gain-optimal, if it is 0-discount-optimal it is also called bias-optimal.




Fig. 1: A MDP with a single action choice in state 1, i.e. two different deterministic policies (Mahadevan, 1996c, adapted).

Furthermore, if a policy is ∞ -discount-optimal then it is said to be Blackwell-optimal (Blackwell, 1962). That is, for Blackwell-optimal policies π^* there exists a discount factor $\gamma^* < 1$ such that $V_{\gamma^*}^{\pi^*}(s) \geq V_{\gamma}^{\pi}(s)$ for all $\gamma \geq \gamma^*$ and under all policies π (Blackwell, 1962; Mahadevan, 1996d). Informally that means there exists a discount factor $\gamma < 1$ which finds Blackwell-optimal policies. However, (i) in more complex, i.e. real world, MDPs this value can be arbitrary close to 1 and (ii) the difference in state values may be very small, which (iii) due to the need of state value function approximation likely causes errors when choosing among different actions in the discounted framework. Schneckenreither (2020) establishes a practically relevant definition for near-Blackwell-optimality.

Definition 2. *If an algorithm infers for any MDP bias-optimal policies, and for a given MDP can in theory be configured to infer Blackwell-optimal policies, but in practise this ability is naturally limited due to the finite accuracy of floating-point representation of modern computer systems, it is said to be near-Blackwell-optimal. An according to a near-Blackwell-optimal algorithm inferred Blackwell-optimal policy is called near-Blackwell-optimal.*

Note that standard discounted RL does not meet the requirements of near-Blackwell-optimality, as it generally does not infer bias-optimal policies. Only by chance the resulting policy could be bias-optimal. The corresponding Bellman equation defined as $V_{\gamma}^{\pi}(s) = \mathbb{E}_{\pi}[r + \gamma V_{\gamma}^{\pi}(s')]$ (see e.g. Sutton, Barto et al., 1998, p.70) provides a way to assess the state values using two consecutively observed states s, s' and the observed reward r returned by the system. When this formula is used as an update rule it provides an algorithm that converges the state values by iteratively adapting the state value estimates. There are four major issues with discounted RL, which results in the fact that it is inapplicable for many problems of operations research.

1. In general standard discounted RL can only infer suboptimal policies as the discount factor is strictly less than one, i.e. $\gamma < 1$.
2. It is very difficult to specify a desired balance between the short-term and long-term (average) rewards. This results from the fact that the long-term

rewards are scaled exponentially, where the factor depends on the chosen γ value.

3. Episodic MDPs with an average reward per step that is non-zero cannot be solved correctly, as the average reward is ignored in the terminal states.
4. The average reward is independently assessed for each state, even though it is the dominating part for all state values and thus is shared between more than one or, for unichain MDPs, even all states. This usually leads to an exponentially increased number of learning steps required for policies to converge.

Especially due to the fourth issue standard discounted RL is not well applicable in the operations research domain. This mainly results from the fact that a continuous assessing of the state values using costs or profit is often required, which leads to intractable long learning phases with high exploration rate even for small sized problems. But as high exploration rates produce errors in the state value estimations (Miller and Veinott, 1969) finding a well working parameterisation is difficult. Average reward RL overcomes these issues by separately assessing the subterms of the state values directly.

2.5 Average Reward Reinforcement Learning

Due to Howard (1964) the average reward $\rho^\pi(s)$ of a policy π and a starting state s is defined as

$$\rho^\pi(s) = \lim_{N \rightarrow \infty} \frac{\mathbb{E}[\sum_{t=0}^{N-1} R_t^\pi(s)]}{N}.$$

In the common case of unichain MDPs, in which only a single set of recurrent states exists, the average reward $\rho^\pi(s)$ is equal for all states s (Mahadevan, 1996b; Puterman, 1994). In the sequel we focus on unichain MDPs in this work and thus may simply refer to it as ρ^π . A policy π^* that maximises the average reward $\rho^{\pi^*}(s) - \rho^\pi(s) \geq 0$ in every state s as compared to any other policy π is called gain-optimal. Mahadevan (1996a) shows that n -discount-optimality for $n = -1$ describes gain-optimality.

Further, for an unichain aperiodic³ MDP problem, such as ergodic MDPs are, the average adjusted sum of rewards or bias value is defined as

$$V^\pi(s) = \lim_{N \rightarrow \infty} \mathbb{E}\left[\sum_{t=0}^{N-1} (R_t^\pi(s) - \rho^\pi)\right],$$

where again $R_t^\pi(s)$ is the reward received at time t , starting in state s and following policy π . Note that the bias values are bounded due to the subtraction of the average reward. Thus the bias value can be seen as the rewards that

³ In the periodic case the Cesaro limit of degree 1 is required to ensure stationary state transition probabilities and thus stationary bias values (Puterman, 1994). Therefore to ease readability we concentrate on unichain and aperiodic MDPs.




Fig. 2: A MDP with two gain-optimal deterministic policies, π_A going left in 1 is also bias-optimal, while π_B is not (Adapted from Mahadevan, 1996b).

additionally sum up in case the process starts in state s . A policy π^* that is gain-optimal is also bias-optimal, if it maximises the bias values $V^{\pi^*}(s) - V^\pi(s) \geq 0$ in every state s and compared to every other policy π . Bias-optimality is given by setting $n = 0$ in the definition of n -discount-optimality (Mahadevan, 1996a).

Especially for highly probabilistic systems bias-optimality is important. To clarify this consider Figure 2 which again consists of two possible deterministic policies with the only choice in state 1. Both policies have the same average reward of 1. However, only taking the A loop is bias-optimal, as its policy π_A leads to bias values $V^{\pi_A}(1) = 0.5$, $V^{\pi_A}(0) = -0.5$ and $V^{\pi_A}(2) = 1.5$, while policy π_B which selects the B loop generates bias values $V^{\pi_B}(1) = -0.5$, $V^{\pi_B}(0) = -1.5$ and $V^{\pi_B}(2) = 0.5$. This yields from the fact that the actions with non-zero rewards are selected earlier in policy π_A . Consider starting in state 1. Under the policy π_A the reward sequence is $(2, 0, 2, 0, \dots)$, while for the other policy π_B it is $(0, 2, 0, 2, \dots)$.

2.6 The Laurent Series Expansion of Discounted State Values

Miller and Veinott (1969) established the link between discounted RL state values $V_\gamma^\pi(s)$ and average reward RL values $\rho^\pi(s)$ and $V^\pi(s)$ using the Laurent series expansion as

$$V_\gamma^\pi(s) = \frac{\rho^\pi(s)}{1 - \gamma} + V^\pi(s) + e_\gamma^\pi(s),$$

where the error term $e_\gamma^\pi(s)$ exists of infinitely many subterms and Puterman (1994) shows that $\lim_{\gamma \rightarrow 1} e_\gamma^\pi(s) = 0$. Note how the first term depending on the average reward $\rho^\pi(s)$ converges to infinity as γ increases.

However, an important insight is the connection between n -discount-optimality and the Laurent series expansion of $V_\gamma^\pi(s)$. Each addend corresponds to one step in n -discount-optimality. That is for (-1) -discount-optimality the addend describes (a scaled version of) the term to maximise for gain-optimality, for 0-discount-optimality the term to maximise for bias-optimality, and finally n -discount-optimality for $n \geq 1$ requires to maximise the error term $e_\gamma^\pi(s)$. The later only exists as γ is strictly less than 1. Thus this term incorporates the number of expected steps and their corresponding reward on the path to the Blackwell-optimal policy. Put differently, it optimises the expected reward collected to reach the Blackwell optimal policy according to the occurrence on the

paths, where shorter paths and those which collect the rewards sooner are preferred.

These addends, where $n = -1, 0, \dots$ denote the coefficients of the Laurent series expansion and thus correspond to the n of the definition of n -discount-optimality, can be reformulated to following constraint problem (Miller and Veinott, 1969; Puterman, 1994, p.346).

$$\rho^\pi(s) - \mathbb{E}[\rho^\pi(s)] = 0 \quad \text{for } n = -1 \quad (1)$$

$$\rho^\pi(s) + V^\pi(s) - \mathbb{E}[V^\pi(s)] = R^\pi(s) \quad \text{for } n = 0 \quad (2)$$

$$W_{n-1}^\pi(s) + W_n^\pi(s) - \mathbb{E}[W_n^\pi(s)] = 0 \quad \text{for } n \geq 1, \text{ where } W_0^\pi(s) = V^\pi(s) \quad (3)$$

The error term $e_\gamma^\pi(s)$ is given by $e_\gamma^\pi(s) = \sum_{n=1}^{\infty} W_n^\pi(s)$ and the expected reward $R^\pi(s) = \mathbb{E}_\pi[r(s, a)]$. Puterman (1994, p.343ff) shows that due to the given degree of freedom, if $n = -1, 0, \dots, M$ constraints are satisfying the above conditions for all states s , then only $\rho^\pi(s), V^\pi(s), W_1^\pi, \dots, W_{M-1}^\pi$ are unique, whereas W_M^π is offset by the vector u , where for the transition probability matrix P the vector u is characterised by $(I - P)u = 0$. Note that u is determined by the number of closed irreducible classes of P , that is for ergodic MDPs u is determined by a single constant. Average reward learning is based on the above formulation.

A major problem occurring at average reward RL is that the bias values are not uniquely defined without solving the first set of constraints defined by the error term addends (see Mahadevan, 1996d; Puterman, 1994, p.346). Our average reward adjusted learning algorithm (ARAL), based on the tabular version of Schneckenreither (2020), does not require the exact solution for $V^\pi(s)$, but a solution which is offset suffices. Clearly this observation reduces the required iteration steps tremendously as finding the exact solution, especially for large discount factors, is tedious. Therefore, we allow to set $\gamma = 1$, which induces $X_\gamma^\pi(s) = V^\pi(s) + u$, where u is for unichain MDPs a scalar value independent of s , i.e. equivalent for all states of the MDP (Puterman, 1994, p.346). If we are interested in correct bias values, i.e. γ is sufficiently close but strictly less than 1, our approach is a tremendous advantage over average reward RL as it reduces the number of iterative learning steps by requiring only a single constraint per state plus one for the scalar average reward value. That is, for an MDP with N states only one more constraint ($N + 1$) has to be solved in ARAL as compared to (at least) $2N + 1$ nested constraints for average reward RL. Therefore, it is cheap to compute $X_\gamma^\pi(s)$, while it is rather expensive to find the correct values of $V^\pi(s)$ directly, especially in an iterative manner as RL is.

2.7 Average Reward Adjusted Reinforcement Learning

Thus, average reward adjusted RL can be seen as a specialised version of average reward RL, where it targets unichain MDPs only. In the sequel we briefly present the needed formalism of average reward adjusted RL. For more comprehensive version we refer to Schneckenreither (2020).

Definition 3. The average reward adjusted (discounted) state value $X_\gamma^\pi(s)$ of a state s under policy π and with discount factor $0 \leq \gamma \leq 1$ is given as $X_\gamma^\pi(s) := V^\pi(s) + e_\gamma^\pi(s)$.

With the Laurent Series expansion this reformulates to $X_\gamma^\pi(s) = V_\gamma^\pi(s) - \frac{\rho^\pi}{1-\gamma}$. Therefore, $X_\gamma^\pi(s)$ describes the discounted state value adjusted by the average reward term. The Bellman Equation is given by $X_\gamma^\pi(s) = \mathbb{E}_\pi[r_t + \gamma X_\gamma^\pi(s_{t+1}) - \rho^\pi \mid s_t = s]$. The corresponding Bellman optimality equation then is $X_\gamma^{\pi^*}(s) = \max_a \mathbb{E}_{\pi^*}[r_t + \gamma X_\gamma^{\pi^*}(s_{t+1}) - \rho^\pi \mid s_t = s]$. Turning this formula into an update rule it builds the foundation of our algorithm.

3 Average Reward Adjusted Deep RL Algorithm

This section presents the average reward adjusted deep RL algorithm (ARAL). First, we present a basic version, which suffices for small problem sizes and introduce the idea of independent agents, before presenting an scalable optimised version.

3.1 Basic Average Reward Adjusted Deep RL Algorithm

The model-free average reward adjusted RL algorithm is based on the tabular version of Schneckenreither (2020) and depicted in Algorithm 1. The algorithm

Algorithm 1 Basic near-Blackwell-optimal deep RL for unichain MDPs

- 1: Initialise state s_0 and network parameters θ_X^T, θ_X^W randomly, set an exploration rate $0 \leq p_{\text{learn}} < p_{\text{expl}} \leq 1$, exponential smoothing learning rates $0 < \alpha, \gamma < 1$, and discount factors $0 < \gamma_0 < \gamma_1 \leq 1$, where $\gamma_1 = 1$ is usually a good choice.
 - 2: **while** the stopping criterion is not fulfilled **do**
 - 3: With probability p_{expl} choose a random action and probability $1 - p_{\text{expl}}$ one that fulfills $\max_a \preceq_\epsilon (X_{\gamma_1}^\pi(s_t, a; \theta_X^W), X_{\gamma_0}^\pi(s_t, a; \theta_X^W))$. Let a_t^{rd} indicate if the action was chosen randomly.
 - 4: Carry out action a_t , observe reward r_t and resulting state s_{t+1} . Store the experience $(s_t, a_t, a_t^{\text{rd}}, r_t, s_{t+1})$ in the experience replay memory M .
 - 5: Sample random mini-batch of m experiences from M
 - 6: Reset gradients: $d\theta_X^W \leftarrow 0$
 - 7: **for** each experience i with $(s_t, a_t^{\text{rd}}, a_t, r_t, s_{t+1})$ **do**
 - 8: **if** a non-random action was chosen or $p_{\text{expl}} > p_{\text{learn}}$ **then**
 - 9: $\rho^\pi \leftarrow (1 - \alpha)\rho^\pi + \alpha[r_t + \max_a X_{\gamma_1}^\pi(s_{t+1}, a; \theta_X^W) - X_{\gamma_1}^\pi(s_t, a_t; \theta_X^W)]$
 - 10: $y_{i, \gamma_0} \leftarrow r_t + \gamma_0 \max_a X_{\gamma_0}^\pi(s_{t+1}, a; \theta_X^T) - \rho^\pi$
 - 11: $y_{i, \gamma_1} \leftarrow r_t + \gamma_1 \max_a X_{\gamma_1}^\pi(s_{t+1}, a; \theta_X^T) - \rho^\pi$
 - 12: Sum gradients on $(y_{i, \cdot} - X_{\gamma_1}^\pi(s_{t+1}, a; \theta_X^W))^2$ wrt. parameters θ_X^W in $d\theta_X^W$
 - 13: Perform update with average of the gradients $d\theta_X^W/m$ on θ_X^W
 - 14: Every C steps exponentially set target network: $\theta_X^T \leftarrow (1 - \gamma)\theta_X^T + \gamma\theta_X^W$
 - 15: Set $s \leftarrow s'$, $t \leftarrow t + 1$ and decay parameters
-

operates on a target network θ_X^T and a worker network θ_X^W as in deep Q-Learning (Mnih et al., 2015). In the action selection process a ϵ -sensitive lexicographic order $a = (a_1, \dots, a_n) \prec_\epsilon (b_1, \dots, b_n) = b$ is used. It is defined as $a \prec_\epsilon b$ if and only if $|a_j - b_j| \leq \epsilon$ for all $j < i$ and $|a_i - b_i| > \epsilon$. Note that the resulting sets of actions may not be disjoint, but taking the maximum as required in our algorithm is straight-forward and thus cheap to compute. The first action selection criteria selects the actions which maximise the bias values, the second the error term. Equation 2 of the average reward RL constraints is used to estimate the average reward based on the currently inferred state values of two consecutive states and the returned reward. This infers a solid average reward prediction of the current policy in few steps only and thus quickly leads to better policies, as ones with higher gain are preferred. Schneckenreither (2020) and Tadepalli and Ok (1998) find that updating the average reward by Equation 2 is superior as compared to exponentially smoothing the actual observed rewards. This makes sense as the later rather describes policies from the past.

The average reward adjusted state value, which is approximated based on the Bellman optimality equation, is estimated twice. Once to infer the bias values with a very high discount factor γ_1 , e.g. $\gamma_1 = 1.0$, and a second time with a discount factor $\gamma_0 < \gamma_1$. The later is used to more selectively choose actions in case more than one action leads to bias-optimal policies. The predetermined discount value γ_0 is a measure for farsightedness. Lower values prefer to partially collect reward in fewer steps and higher values are used to prefer actions for which the full bias is collected earlier. The corresponding update values are saved to calculate the gradients based on the quadratic errors. Every C steps the target network smoothly adapts the parameters of the worker network.

3.2 Independent Agents

In our application the agent can update the lead time for each product type by a maximum magnitude of 1 in regard to the lead time of the current period. A naive approach is to define one action for each combination of possible lead time updates over all product types, e.g. one action for a system with six product types could be the vector $(+1, +1, +0, +0, +1, +0)$ for which each component defines the update of the corresponding product type lead time from the current lead time. This however, due to the exponential growth, results in only being able to solve very small production systems with only few product types. A common solution in production planning is to aggregate the products into product groups, cf. hierarchical production planning (Schneeweiß, 1995). Although applicable in our setup if we group the products into bottleneck and non-bottleneck products as done by Schneckenreither, Haeussler, and Gerhold (2020), we refrain from this adaption as it does not scale to larger sized production systems.

Therefore, in the seek of a broader solution we establish the idea of *independent agents*. Each independent agent operates on the same system state, but chooses only a part of the actual executed action. In our case each independent agent is responsible for the lead time update of a specific product type. To the best of the author’s knowledge this is the first work to reduce the action space




Fig. 3: Illustration of independent agents as agent one increases the lead time

in this way. Figure 3 illustrates the idea, where wlog. we use two agents and for simplicity omit the reward and truncate the state to the lead times. Assume that the first independent agent chooses to increase the lead time from 3 to 4. Thus, the future state, seen from this agent, depends on the probability of the other independent agent choosing to increase (\uparrow), decrease (\downarrow) or not change ($-$) the lead time. The Markov property, that is $p(s_{t+1}, r_t | s_t, a_t) = p((4, \cdot), r_t | (3, 3), (\uparrow, \cdot))$, is not touched by this adaption. While for the first agent the reward and future state only depends on the current state and action taken, the corresponding probabilities for future state and reward change during the learning process as the agent starts to behave differently.

3.3 Scalable Average Reward Adjusted Deep RL Algorithm

The basic algorithm above suffices only for problem sets of rather small size due to the computational complexity and the growth of the action space. In our case, that would either require a production system with at most 3 product types, or an aggregation of product types into product groups. In the sense of this foundational research paper, we aim for an enhanced version of the algorithm that is scalable.

Therefore, our implementation includes several optimisations, that help the algorithm in better exploring the solution space and stabilise the learning process. These include the above discussed independent agents that operate on the same state, n -step learning (Mnih et al., 2016), and ensuring to constantly reduce the average reward as searching for policies with higher average costs is unnecessary. Algorithm 2 depicts a detailed view on the algorithm including these extensions and with the objective to minimise the returned rewards, as we aim for minimal costs in our experiments. The idea is the same as in Algorithm 1, except that the action selection procedure is a loop over all independent agents and the implementation of the n -step logic. The agent first chooses an action (steps 4 – 5), executes the selected action and stores the experience in the replay memory (step 6) and then samples a mini-batch of experiences each of length $nstep$ to improve the policy (steps 7 – 21). Step 17 uses an exponen-

Algorithm 2 Optimised near-Blackwell-optimal deep RL for unichain MDPs

```
1: Initialise state  $s_0$  and network parameters  $\theta_X^T, \theta_X^W$  randomly, set an exploration
   rate  $0 \leq p_{\text{learn}} \leq p_{\text{expl}} \leq 1$ , exponential smoothing learning rates  $0 < \alpha_{\text{max}}, \alpha, \gamma < 1$ ,
   a sufficiently large default value for  $\rho_{\text{max}}^\pi$  and discount factors  $0 < \gamma_0 < \gamma_1 \leq 1$ 
2: while the stopping criterion is not fulfilled do
3:    $a_t^{\text{rd}} \leftarrow \text{False}$ 
4:   for each independent agent  $i \in \text{ag}_1, \text{ag}_2, \dots, \text{ag}_n$  do
5:     With probability  $\frac{p_{\text{expl}}}{n}$  choose action  $a_{t,i}$  randomly or otherwise one of the
     set defined by  $\min_{a \in \mathcal{A}(i)} \preceq_\epsilon (X_{\gamma_1}^\pi(s_t, a; \theta_X^W), X_{\gamma_0}^\pi(s_t, a; \theta_X^W))$ . Set  $a_t^{\text{rd}} \leftarrow \text{True}$ 
     if the action was chosen randomly.
6:   Carry out action  $a_t = (a_{t,\text{ag}_1}, \dots, a_{t,\text{ag}_n})$ , observe reward  $r_t$  and resulting state
    $s_{t+1}$ . Store experience  $(s_t, a_t, a_t^{\text{rd}}, r_t, s_{t+1})$  in the experience replay memory  $M$ .
7:   if  $t \bmod n_{\text{step}} \equiv 0$  then
8:     Reset gradient:  $d\theta_X^W \leftarrow 0$ 
9:     Sample  $m$  random mini-batches of  $n_{\text{step}}$  consecutive experiences from  $M$ 
     numbered by  $1 \leq i \leq n_{\text{step}}$ , each with  $(s_i, a_i^{\text{rd}}, a_i, r_i, s_{i+1})$ .
10:    for each mini-batch do
11:      for each experience, starting with the latest seen experience  $n_{\text{step}}$  do
12:        if experience  $n_{\text{step}}$  or  $\neg a_i^{\text{rd}}$  or  $p_{\text{expl}} > p_{\text{learn}}$  then
13:           $\text{target}_{\gamma_0} \leftarrow \min_a X_{\gamma_0}^\pi(s_{i+1}, a; \theta_X^T)$ 
14:           $\text{target}_{\gamma_1} \leftarrow \min_a X_{\gamma_1}^\pi(s_{i+1}, a; \theta_X^T)$ 
15:          if  $\neg a_i^{\text{rd}}$  or  $p_{\text{expl}} > p_{\text{learn}}$  then
16:             $\rho^\pi \leftarrow (1 - \alpha)\rho^\pi + \alpha[r_i + \min_a X_{\gamma_1}^\pi(s_{i+1}, a; \theta_X^W) - X_{\gamma_1}^\pi(s_i, a_i; \theta_X^W)]$ 
17:            Ensure improvement of  $\rho^\pi$ , otherwise reset to old value
18:             $\text{target}_{\gamma_0} \leftarrow r_i + \gamma_0 \text{target}_{\gamma_0} - \rho^\pi$ 
19:             $\text{target}_{\gamma_1} \leftarrow r_i + \gamma_1 \text{target}_{\gamma_1} - \rho^\pi$ 
20:             $d\theta_X^W \leftarrow d\theta_X^W + \partial(\text{target}_{\gamma_0} - X_{\gamma_0}^\pi(s_{i+1}, a; \theta_X^W))^2 / \partial \theta_X^W$ 
21:             $d\theta_X^W \leftarrow d\theta_X^W + \partial(\text{target}_{\gamma_1} - X_{\gamma_1}^\pi(s_{i+1}, a; \theta_X^W))^2 / \partial \theta_X^W$ 
22:          Perform update of  $\theta_X^W$  using  $d\theta_X^W / (m \cdot n_{\text{step}})$ 
23:        Every  $C$  steps exponentially set target network:  $\theta_X^T \leftarrow (1 - \gamma)\theta_X^T + \gamma\theta_X^W$ 
24:        Set  $s \leftarrow s', t \leftarrow t + 1$  and decay parameters
```

tially smoothed overestimation (by 2.5%) of the average reward and rate 2^{-5} . All average reward updates above this overestimation are rejected.

As in (Mnih et al., 2016, 2015) the algorithm uses a *Replay Memory* to save experiences and then randomly selects experiences from there for learning. This ensures that the independent and identically distributed (i.i.d.) assumption of the samples for the ANN backpropagation algorithms are satisfied. We extend this idea by splitting the available size of the replay memory into subsets, where each subset represents one action and sample uniformly over these subsets first. This prevents overfitting when one action is overrepresented in the policy, e.g. no lead time change occurs much more often than the other actions.

N-step learning collects a set of consecutive experiences, in our case retrieved from the replay memory, and connects the state values starting from the latest observed experience, such that the newly computed state value of the corre-

sponding experience propagates to the experience of the previous period. This allows a faster convergence, as values bootstrap over multiple periods in contrast to just one step. The experiences are stored according to the first taken action in the n -step series. We uniformly sample over the actions to ensure each action is constantly trained. This adaption reduces catastrophic forgetting of already accumulated knowledge. Furthermore, we use 8 additional **worker threads** (parallel RL) that operate on the same environment to collect a broader range of experiences (as e.g. in Mnih et al., 2016). The average reward value is shared and computed as average over the main agent and all workers.

The algorithm initiation strategy is to first fill the replay memory buffer by skipping any learning (steps 7 – 21 in Algorithm 2) until the period, that corresponds to the replay memory size, is reached. Then for 500 steps the average reward is set to the last computed exponentially smoothed value to prevent divergence due to the untrained network weight initialisation values.

4 Experimental Evaluation

This section introduces the simulation model and the parameterisation of the algorithm. To provide a proof-of-concept for the proposed algorithm we adapt the flow shop presented in Schneckenreither, Haeussler, and Gerhold (2020). However, our algorithm operates on discrete lead times on a product type basis and thus is not directly comparable to their order based lead time management approach.

4.1 Simulation Model

We adapt the simulation model of a make-to-order (MTO) flow shop with six different products and six work-centers of Schneckenreither, Haeussler, and Gerhold (2020) and examine the system under moderate and high utilisation in combination with moderate or high variability in demand and processing times. This section provides the details of the simulation model and the experimental setup.

The simulation model is depicted in Figure 4, where the customers initiate the process by placing orders as indicated by the dashed line. The manufacturer produces six different product types, each of which is identified by a natural number $1, 2, \dots, 6$ and has a different routing through the production system. The routing setup is provided by corresponding labels of the edges. For instance, product types 1 – 3 are routed from machine $M1$ to machine $M2$, while the other products are transferred to machine $M3$. The production system consists of three production stages with diverging material flow and no return visits. The incoming orders are uniformly distributed among the product types. A period lasts 960 minutes (= 16 hours) which represent two 8-hour shifts. At the beginning of each period orders can be released into the production system assuming material availability for all orders at their release date. Upon release, the orders are placed in the buffer at machine $M1$. Buffers and inventories are




Fig. 4: Production System of the Simulation Model with routing, processing time distributions and demand interarrival time distributions.

plotted as filled squares. Thus, each workstation is equipped with a buffer which queues the orders until processing is started. All buffers use a first-come-first-serve dispatching rule. Furthermore, each workstation processes only one order at a time and early deliveries are prohibited. Orders completed ahead of their due date remain in the finished goods inventory (FGI) until they are due and sent to the customer.

Processing Times. The machine processing time distributions are specified under the corresponding node labels of the machines. To simulate high variance the processing times for all machines are drawn from an exponential distribution. There is one bottleneck machine ($M5$) and therefore we refer to product types 2 and 5, those routed through machine $M5$, as bottleneck products whereas products 1, 3, 4, 6 are non-bottleneck products.

Demand. The incoming orders are placed in the order pool with a due date slack of 7 periods, that is each order arriving within period t will be due at the end of period $t + 7$. The due date slack of 7 was chosen to (i) ensure sufficient time between the first occurrence of orders in the order pool and the latest possible release of orders and (ii) to have a clear cause and effect relationship between the order release decision and the cost performance in the analysis. To simulate medium and high variability of the demand process the interarrival time between consecutive incoming orders is either drawn from an exponential (Exp) or an uniform (\mathcal{U}) distribution, which are adjusted to yield the desired bottleneck utilisation level of either 70% with Exp(135) and $\mathcal{U}(95, 175)$, 80% with Exp(118) and $\mathcal{U}(78, 158)$, or 90% with Exp(105) and $\mathcal{U}(65, 145)$. These values were chosen as they comprise the non-linear relationship between flow times and high utilisation levels.

Order Release. Before the order release decision is made, all orders in the order pool are sorted by due date. This portrays the implementation of a sequenc-

ing rule. According to this sequence, orders are considered for release in the beginning of each period starting with the highest priority order.

In our model, orders are released by specifying lead times $LT_1, LT_2, \dots, LT_6 \geq 1$, where the index corresponds to the product type and $LT_i \in \mathbb{N}$. By setting a lead time LT_i a *planned release date* is computed for each order j of product type i in the order pool given by

$$PRD_j = DD_j - LT_i, \quad (4)$$

where DD_j denotes the due date of the order. For dynamic order releases the lead times LT_i may vary from period to period, whereas in static order release methods the lead times are predetermined and fixed (Kim and Bobrowski 1995b; Ragatz and Mabert 1988). For both approaches a job of product type i with lead time LT_i in period t and due date DD_j is released at the end of period t if and only if $t \geq PRD_j$. Thus, in the dynamic setting the planned release date PRD_j of order j can be updated several times before it is actually released to the production system. However, once an order is released its release cannot be revoked. Hence, if the set lead time corresponds with the actual flow time the product is finished in the same period as it is shipped. If the lead times are set either too long or short, the order has to wait in the FGI until it is due or the order is late and backorder (BO) costs occur.

Costs. The cost parameters are set by assuming an increase in value from raw material (1 Dollar per order and period) to the final product (4 Dollar per order and period) and the backorder costs are set very high (16 Dollar per order and period) due to the MTO environment. All costs are assessed based on the production system state at the end of each period.

5 Lead Time Setting

In this section we define the tested lead time setting algorithms. This includes a description of the MDP used for the reinforcement learning agents.

5.1 Conventional Order Release Rule

As external benchmark for comparison we use different parameterized backward infinite loading (BIL) techniques (Ackerman, 1963). BIL uses Equation 4 with a predetermined and fixed lead times LT_i .

5.2 Average Reward Adjusted Deep RL Algorithm

Markov Decision Process. The underlying MDP is unichain and looks as follows. Both, state space \mathcal{S} and action space \mathcal{A} are discrete. For the *state space* we adapt the result of Knollmann and Windt (2013), which state that the lead times ought to control the WIP level, instead of the due date reliability as changes are directly seen. Thus, any state $s \in \mathcal{S}$ of the state space is composed of the following information for each product type pt :

- The currently set lead time $LT_g \in \{1, 2, \dots, dds\}$ (Recall: Due Date Period – Lead Time = Release Period). Note that we bound the maximum lead time with due date slack dds , which is 7 in our setup.
- Counters $OP_{g,d} \in \mathbb{N}$ for the number of orders in the order pool divided in time buckets with $d \in \{1, 2, \dots, dds\}$, which stands for the number of periods until the due date.
- Counters $Q_{g,i,d} \in \mathbb{N}$ for the number of orders in the machine and queue $i \in \{M_1, M_2, \dots, M_6\}$ divided in time buckets with $d \in \{1, 2, \dots, dds\}$, which stands for the number of periods until the due date.
- Counters $FGI_{g,d} \in \mathbb{N}$ of orders in the finished goods inventory divided in time buckets with $d \in \{1, 2, \dots, dds\}$, which stands for the number of periods until the due date. Orders with a due date with more than 3 periods ago are listed in the counter FGI_{-3} .
- Counters $S_{g,d} \in \mathbb{N}$ of shipped orders from the last period divided in time buckets with $d \in \{-3, -2, \dots, dds\}$, which stands for the number of periods until the due date. Orders with a due date with more than 3 periods ago are listed in the counter S_{-3} .

The algorithm implicitly learns a function which maps the current state of the production system to a lead time update and thus a release decision. In an optimal situation it does so by multiply exploring every action for each state and assessing its economic viability. Orders are released once the due date is within the interval $[t, t + \text{lead time}]$, whereas t is the current period. Clearly, this yields bulk releases, meaning that either all or no orders with the same due date and of the same product type are released.

The *action space* is composed of two decisions, the relative change of the lead times to the currently set lead times $LT_{pt} \in \{1, 2, \dots, dds\}$ for each product type pt . We restrict the lead time update for state s_{t+1} according to the lead time LT_{pt} from state s_t by a maximum change of 1. Thus, if LT_{pt} is the current lead time of product type pt the agent can choose a lead time in $\{1, 2, \dots, dds\} \cap \{LT_{pt} - 1, LT_{pt}, LT_{pt} + 1\}$. Recall that we implement independent agents to reduce the action space. In particular, this adaption enables us to present the algorithm the above defined state space not only on the basis of product groups, but directly for each product type. Note that without this extension the problem is computationally infeasible due to $3^6 = 729$ possible actions. Removing the dependency between the lead times of different product types reduces the set of possible actions, and thus the number of ANN output layer nodes, to $3 \cdot 6 = 18$. In preliminary experiments on small examples we only observed a very slight increase of required steps during training. Therefore, this is a vital adaption to ensure scalability.

Reward. In each period the agents choose actions, which when executed generate a reward while traversing to the next period by simulating the production system. The rewards are the accumulated costs at the end of the period consisting of the backorders, the WIP orders and the number of orders in the FGI. The objective is to minimise the returned reward.

Neural Network Setup. Due to the vast possibilities the neural network architecture was manually optimised, starting with less stochastic systems and aiming for a coffin shape that can be used in all setups. During this process we constantly adapted and reevaluated the best network on various production system configurations. The resulting architecture depends on the number of features n_{inp} and the number of output nodes n_{out} , which itself depends on the number of actions $n_{\text{act}} = |\mathcal{S}|$, the independent agent configuration, and the number of functions that are approximated. For the later, in case of ARAL this is 2, as we represent both state-action value functions $X_{\gamma_0}^\pi$ and $X_{\gamma_1}^\pi$ with the same neural network, while deep Q-Learning (cf. Mnih et al., 2015) only approximates the function Q_γ^π . The resulting feedforward deep ANN architecture is given as $n_{\text{inp}} \rightarrow 1.75n_{\text{inp}} \rightarrow 1.5n_{\text{inp}} \rightarrow n_{\text{out}}$, where each arrow symbolises a fully-connected layer with a leaky rectified linear unit (ReLU) as activation function, which leaks values smaller than 0 with rate 0.02. Note that we explicitly do not use an output activation function. The idea is that we do not want to restrict the approximations to values inside this range. This is especially important for Q-Learning, as the range of the state values are difficult to forecast. The actual values n_{inp} and n_{out} depend on the used algorithm. In our case using 6 independent agents we gain $n_{\text{inp}} = 324$ and $n_{\text{out}} = 36$.

The features representing order counts are scaled using min-max scaling with minimum 0 and maximum 5, the lead times with minimum 0 and maximum 7. To further ensure smaller ANN parameter values the output is transformed using min-max scaling to the interval $(-400, 800)$, s.t. an ANN output of 1 represents a value of 800. On each training step, i.e. every n -step steps, randomly m sets of experiences the training mini-batch size each consisting of n -step experiences are selected, there gradients computed and the sum of gradients backpropagated through the worker network using Adam (Kingma and Ba, 2014), where the weight decay with parameter λ is implemented as presented by Loshchilov and Hutter (2017). Every C steps the target network is updated by exponentially smoothing the parameters with rate γ . As neural network initialisation method we use the commonly known Xavier initialisation (Glorot and Bengio, 2010), i.e. the initial value for each weight i is sampled from a uniform distribution $W_{l,i} \sim \mathcal{U}(-1/\sqrt{n_l}, 1/\sqrt{n_l})$, where n_l is the number of nodes in layer l .

5.3 Deep Q-Learning Algorithm

Furthermore, we adapt the deep Q-Learning algorithm of Mnih et al. (2015) with integrated extensions as described above, i.e. the adaptations to n -step learning, multiple workers, independent agents, scaling of the output, smoothly adapting the target network. The resulting Algorithm 3 is provided in the Appendix.

To ensure comparability we use the same neural network setup as for ARAL, except that we halve the number of output nodes, as the Q-Learning algorithm only approximates one state-action value function. Therefore the output nodes n_{out} are 18, while the input nodes $n_{\text{inp}} = 324$ stay the same.

Parameter	α	γ_0	γ_1	γ	ϵ	p_{exp}	p_{learn}
Start value	0.01	0.8	0.99/1.0	0.01	0.25	1.0	0.5
Exp. Smth. Rate	0.25			0.25		0.25	
Exp. Smth. Steps in 10^3	50			100		100	
Minimum value	5^{-5}			0.001		0.005	
<hr/>							
Configuration	Value						
Workers Min. Exploration	0.01, 0.02, 0.03, 0.04, 0.05, 0.10, 0.15, 0.20						
Target Network Update	Exp. smooth. with rate γ every $C = 100$ steps						
N-Step	5						
Training Mini-Batch Size	4						
Adam ANN Backpropagation	$\alpha = 0.005, \beta_1 = 0.9, \beta_2 = 0.999, \epsilon = 10^{-8}, \lambda = 0.001$						
Replay Memory Size	10800						

Table 1: Overview of Parameter Setup.

6 Computational Experiments

This section presents the experimental results and gives an overview of the performance of the established algorithm compared to the conventional static lead time setting algorithms.

Table 1 presents an overview of the parameter setup. The parameters α , γ and p_{exp} are exponentially decayed as indicated. Due to the vast possibilities we manually optimised the values. The methodology was to investigate the values on very small production system with just one machine and product due to the computational complexity of bigger sized problems. The resulting values were copied to the problem at hand and further manually optimised in pre-experimental runs. The actual used discount factors γ_1 differ for each method and thus are further specified in the result tables. In a similar approach we reached the configuration values, where a further increase of the values was hindered by the available computational resources.

For the experimental runs we used personal computers with Intel CPUs i7-7700 (16GB RAM) and i9-7900X (32GB RAM). For each method we performed 500.000 learning steps, which took about 2-3 days. For the same experimental setup the inputs/random variables are the same for each period. After learning, each order release method was evaluated on 20 demand streams, each with 1000 periods of warm-up and 7000 periods of evaluation phase. Welch’s procedure was applied to approximate the length of the warm-up period (see Law and Kelton (2000)).

Medium Demand Variability. The results for medium variability in the system are presented in Table 2, where a uniform demand interarrival time is used. We use the tuples 70- \mathcal{U} , 80- \mathcal{U} , and 90- \mathcal{U} to indicate the experimental setup of using uniform demand interarrival times with a aimed bottleneck utilisation of 70%, 80%, or 90% respectively. The results state the accumulated costs, averaged over

70- \mathcal{U}	WIPC	FGIC	BOC	FGI+BO	SUM	SFTT	FGIT	SL(%)	Cost.p.P	ρ^π
QL _{0.99}	38.56	134.70	81.70	216.41	254.97 ¹	1.34	1.22	90.6	36.42	
ARAL _{0.99}	37.20	113.10	100.40	213.50	250.70	1.31	1.11	88.9	35.81	35.04
ARAL _{1.00}	38.66	156.66	59.74	216.40	255.06 ¹	1.34	1.34	93.2	36.44	36.50
BIL1	32.06	0.00	512.90	512.90	544.96	1.20	0.49	44.9	77.85	
BIL2	32.06	83.46	105.01	188.47	220.53	1.20	0.94	88.6	31.50	
BIL3	32.16	247.73	21.68	269.41	301.56	1.20	1.83	97.8	43.08	
BIL4	32.14	429.10	5.08	434.17	466.32	1.20	2.81	99.5	66.62	
80- \mathcal{U}	WIPC	FGIC	BOC	FGI+BO	SUM	SFTT	FGIT	SL(%)	Cost.p.P	ρ^π
QL _{0.99}	63.10	180.40	121.53	301.93	365.03	1.68	1.34	89.7	52.15	
ARAL _{0.99}	65.31	168.49	148.59	317.08	382.39	1.72	1.28	86.4	54.63	52.92
ARAL _{1.00}	60.78	234.05	76.17	310.21	370.99 ¹	1.64	1.59	93.4	53.00	52.29
BIL1	51.89	0.00	830.22	830.22	882.11	1.47	0.49	32.8	126.02	
BIL2	52.00	70.11	255.17	325.28	377.27	1.48	0.82	79.4	53.90	
BIL3	51.98	240.18	77.82	318.00	369.97 ¹	1.48	1.62	94.0	52.85	
BIL4	51.83	442.31	26.71	469.02	520.85	1.47	2.56	98.1	74.41	
90- \mathcal{U}	WIPC	FGIC	BOC	FGI+BO	SUM	SFTT	FGIT	SL(%)	Cost.p.P	ρ^π
QL _{0.99}	98.91	190.68	297.33	488.01	586.92 ¹	2.13	1.28	82.7	83.85	
ARAL _{0.99}	105.35	263.95	207.18	471.13	576.48 ¹	2.24	1.59	87.9	82.35	79.83
ARAL _{1.00}	107.87	240.80	253.87	494.66	602.54 ²	2.28	1.49	84.7	86.08	80.43
BIL1	88.66	0.00	1,418.53	1,418.53	1,507.19	1.96	0.50	21.0	215.31	
BIL2	88.32	50.85	646.00	696.86	785.17	1.96	0.71	64.5	112.17	
BIL3	88.25	207.21	299.89	507.10	595.35 ²	1.96	1.35	85.2	85.05	
BIL4	87.77	415.18	153.61	568.79	656.57	1.95	2.21	92.9	93.80	

Table 2: Results with moderate variability in the system due to uniformly distributed demand interarrival times. ^{1,2} Costs marked with the same number are not significantly different from each other.

all replications, for the work in process (WIPC), the finished goods inventory (FGIC), the backorder costs (BOC), as well as the sum of the finished good inventory and backorder costs (FGI+BO) and the sum over all three cost centers (SUM). To ease readability all costs are given in 1000 Dollars. For a better insight in the timing performance we also provide the average shopfloor throughput time (SFTT), i.e. the average time an order spends in the system after being released and until reaching the finished goods inventory, the finished goods inventory time (FGIT), as well as the service level (SL). The later is given in percentage and states the percentage of orders that were produced on time. Finally, the actual measured costs per period (Cost.p.P) and the by the algorithm learned average reward ρ are stated. The tested methods are the deep Q-Learning algorithm with discount factor 0.99 (QL_{0.99}), the ARAL algorithm with discount factors 0.99 (ARAL_{0.99}) and 1.00 (ARAL_{1.00}), as well as the backward infinite loading

technique BIL with fixed lead times between 1 and 4 as indicated by the number (BIL1-BIL4).

For the case of medium variability with a utilisation of 70% 70-(\mathcal{U}) we can see that BIL2 performs best in the sense of accumulated costs (220.53). The RL algorithm follow, where ARAL_{0.99} is slightly but significantly better then the other two variants (250.70, 254.97, 255.06). The other BIL variants accumulate much more costs. For the timing one can see that ARAL_{1.00} performs very good with a service level of 93.2%. Furthermore, the learned average reward ρ and actual average reward are very similar for both ARAL variants. However, with 70% of utilisation the overhead of the reinforcement learning variants are too big to be able to compete with the simple rule based release mechanism of BIL2 as for changing workload the utilisation changes rather linearly due to the low variability in the system.

In the case of 80- \mathcal{U} the deep Q-Learning variant QL_{0.99} performs best in the cost related measures (365.03 for SUM), followed by BIL3 (369.97) and ARAL_{1.00} (370.99), which are not significantly different from each other. One can see the BIL2 (377.27) performs slightly worse than ARAL_{1.00} and is followed by ARAL_{0.99} (382.39). Again ARAL_{1.00} achieves a very high service level of 93.4% with reasonable sum of cost, which is also reflected in the low Backorder costs. Furthermore, the measured and learned average reward values are quite close. Thus, the results already provide the insight that ARAL performs quite well in terms of timing for the completion of orders. This pays off when confronted with higher probabilistic systems.

For instance, with 90% utilisation and medium variable demand stream, i.e. 90-(\mathcal{U}), we can see that ARAL_{0.99} performs best (576.48), followed by QL_{0.99} (586.92) and BIL3 (595.35), as well as ARAL_{1.00} (602.54), where the later two are not significantly different. In this case ARAL_{0.99} also achieves the best service level of 87.9% within this group of best performing algorithms. The WIPC costs, as well as the SFTT, show that the RL agents have a higher work in process, which is clear as they try to smooth the workload of demand stream. When looking at the timing cost it becomes even more apparent that the agents perform well in timing as ARAL_{0.99} performs much better (471.13) than the other two agents (488.01, 494.66) and all the BIL techniques accumulated more FGI+BO costs (507.10 for the best performing BIL, BIL3).

High Demand Variability. For results of the experiments using exponential demand interarrival times, presented in Table 3 we can see a similar pattern as before. In this case we use the tuples 70-Exp, 80-Exp, and 90-Exp to indicate the corresponding experimental setups.

For 70% utilization of the bottleneck, i.e. 80-Exp, BIL3 performs best with costs of 345.80, followed by ARAL_{0.99} with 358.41, and then by QL_{0.99} and BIL2 (360.70, 364.18). In terms of timing performance, again ARAL outperforms the opponents with FGI+BO costs of 296.99, and BIL3 with the second highest costs (297.17). However, the service level of BIL3 is higher with 93.7% due to the lower shop floor throughput time (SFTT, 1.48 periods for BIL3 versus 1.73 periods for ARAL_{0.99}).

70-Exp	WIPC	FGIC	BOC	FGI+BO	SUM	SFTT	FGIT	SL(%)	Cost.p.P	ρ^π
QL _{0.99}	53.87	142.40	164.43	306.83	360.70 ¹	1.59	1.21	83.8	51.53	
ARAL _{0.99}	61.42	181.70	115.29	296.99	358.41	1.73	1.41	89.3	51.20	50.79
ARAL _{1.00}	56.58	199.31	115.71	315.02	371.60	1.64	1.50	88.7	53.09	50.58
BIL1	48.55	0.00	776.75	776.75	825.30	1.48	0.50	33.7	117.90	
BIL2	48.62	66.93	248.64	315.57	364.18 ¹	1.48	0.83	78.0	52.03	
BIL3	48.63	222.60	74.57	297.17	345.80	1.48	1.61	93.7	49.40	
BIL4	48.51	409.53	23.41	432.94	481.45	1.48	2.55	98.1	68.78	
80-Exp	WIPC	FGIC	BOC	FGI+BO	SUM	SFTT	FGIT	SL(%)	Cost.p.P	ρ^π
QL _{0.99}	87.71	227.49	225.15	452.64	540.34 ¹	2.04	1.50	83.4	77.19	
ARAL _{0.99}	89.84	253.48	185.69	439.16	529.00 ^{1,2}	2.08	1.61	87.6	75.57	70.53
ARAL _{1.00}	87.17	242.92	198.46	441.37	528.54 ^{2,3}	2.03	1.57	86.6	75.51	72.94
BIL1	78.73	0.00	1,259.71	1,259.71	1,338.44	1.88	0.50	22.9	191.21	
BIL2	78.20	52.30	548.74	601.04	679.24	1.88	0.73	65.0	97.03	
BIL3	79.03	200.27	242.10	442.37	521.40³	1.89	1.38	85.8	74.49	
BIL4	79.38	395.49	116.94	512.44	591.82	1.90	2.23	93.8	84.55	
90-Exp	WIPC	FGIC	BOC	FGI+BO	SUM	SFTT	FGIT	SL(%)	Cost.p.P	ρ^π
QL _{0.99}	168.66	222.15	788.66	1,010.81	1,179.46 ¹	3.13	1.37	71.5	168.49	
ARAL _{0.99}	158.03	212.73	701.30	914.03	1,072.06 ²	2.97	1.33	71.9	153.15	115.80
ARAL _{1.00}	160.45	286.07	539.80	825.87	986.32	3.01	1.62	79.6	140.90	109.91
BIL1	146.84	0.00	2,349.37	2,349.37	2,496.21	2.80	0.50	13.3	356.60	
BIL2	143.76	34.11	1,412.75	1,446.86	1,590.62	2.75	0.63	47.1	227.23	
BIL3	144.28	154.15	877.45	1,031.60	1,175.88 ¹	2.76	1.10	70.6	167.98	
BIL4	147.08	335.07	622.14	957.21	1,104.29 ²	2.80	1.81	82.5	157.76	

Table 3: Results with high variability in the system due to exponentially distributed interarrival times of the demand. ^{1,2,3} Costs with the same number are not significant different from each other.

In case of 80-Exp BIL3 performs best with costs of 521.40, but is not significantly better than ARAL_{1.00} with 528.54. Also ARAL_{0.99} performs quite well (529.00). The service level of all three algorithm are good, but the order is reversed: ARAL_{0.99} with 87.6%, ARAL_{1.00} with 86.6% and BIL3 with 85.8%. Again the SFTT of the RL agents are slightly higher, i.e. 0.13-0.20 periods, than the ones of the BIL release procedures. The other three BIL methods, BIL1, BIL2 and BIL4, are performing rather bad (1,338.44, 679.24 and 591.82 respectively). Interesting to see is that in this scenario ARAL_{0.99} and ARAL_{1.00} handle the due date performance better than QL_{0.99}, as can be seen in the backorder costs (BOC). Here the former two accumulate costs of 185.69 and 198.46, while QL_{0.99} faces costs of 225.15. This results in timing costs (FGI+BO) of 439.16 for ARAL_{0.99} and 441.37 for ARAL_{1.00}, which also means a better timing cost than BIL3 with 442.37. For this scenario the actual measured average reward (average costs) of

about 75.5 for both ARAL variants and learned average rewards of 70.53 for ARAL_{0.99} and 72.94 for ARAL_{1.00} are a little bit offset. This probably results from the higher variability and utilisation in the system.

Finally, for the case of 90-Exp ARAL_{1.00} performs much better than the other variants with costs of 986.32. It is followed by ARAL_{0.99} with costs of 1,072.06 and rather far, but not significantly, afterwards follows BIL4 (1,104.29). As before the shopfloor throughput time (SFTT) is higher for the RL variants, and with that of course also the work in process costs (WIPC). The backorder costs of ARAL_{1.00} (539.80) are the lowest of all tested methods. Like for the overall costs also in terms of timing costs (FGI+BO) ARAL_{1.00} performs best (825.87), followed by ARAL_{0.99} (914.03) and BIL4 (957.21). BIL3 and QL_{0.99} perform quite similar in the sense of accumulated costs, 1,175.88 and 1,179.46 respectively. But the other measures show that QL_{0.99} has a better timing performance, thus also a slightly high service level with 71.5% as compared to BIL3 with 70.6%. The best service level is achieved by BIL4 with 82.5%, before ARAL_{1.00} 79.6%, but with a high cost as the difference of SUM is 117.97. Interesting to see is that the measured average reward and the learned average reward are very different for both ARAL variants (about 30-40). This is probably due to the high variability and utilisation in the system, but might also be an indicator that the agents need more time for learning. However, due to the high computational complexity we encountered we leave this to future work.

Overall we can see that the RL agents, and especially the established ARAL algorithm performs very well in systems with high utilisation. In case of the less utilised systems the results reveal that there is still potential for the improvement of RL variants.

7 Conclusion

This paper is one of the first papers that use RL for setting lead times to release orders into the production system. We present a novel average reward adjusted reinforcement learning algorithm that uses an artificial neural network to approximate the state-action value functions. We add several extensions to this algorithm, including n-step learning, and a special experience replay memory that stores its experiences according to the actions taken to prevent catastrophic forgetting. Additionally we invent *independent agents* to reduce the action space by independently choosing a subset of possible actions, while operating on the same state. Additionally we adapt the deep Q-Learning algorithm of Mnih et al. (2015) to be able to present a fair comparison between the algorithms, where we also use several backward infinite loading techniques.

The results suggest that for high utilisation the RL learning agents, and here especially the novel ARAL variant, performs very well. In case of medium utilisation the achieved results are as good as with the BIL variants. And for low utilisation one can see that the RL agents still have room for improvement, as BIL outperforms them. When comparing deep Q-Learning to the established ARAL

algorithm, we can see that for all cases, except one (80- \mathcal{U}), ARAL performs better.

Therefore, in the future we plan to develop ARAL further to a actor-critic state-of-the-art RL algorithm. We hypothesise that due to the more correct state-action value estimations of ARAL as compared to deep Q-Learning, cf. Schneck-enreither (2020), the critic will provide a better feedback of the value function update to the actor. This should lead to better policies than with standard actor-critic algorithms, e.g. A3C (Mnih et al., 2016). Furthermore, one limitation of this proof-of-concept study is that due to the computational complexity we restricted our scope to flow-shop production systems only. In subsequent studies we plan to test the algorithms on job-shop queuing systems as well. Additionally, we want to exploit hierarchical production planning (Schneeweiß, 1995) by implementing agents on different hierarchical levels which can bargain with each other. For instance, one agent might be responsible for medium term planning order release decision at the top level, while another agent concentrates on short-term scheduling at the base level. At the start of the process a proposed order release plan is generated by the first agent. Accordingly the agent from the base level performs a preliminary scheduling of these orders and may report very expensive orders as feedback. This information is then taken into consideration on the top level which adapts the order release plan accordingly. Here the top level agent minimises the overall costs by load leveling on an aggregated model, while at the bottom level the agent (mainly) minimises backorder costs. Finally, although the introduced actions space reduction method, named *independent agents*, works very well, it remains to be shown that the achievable policies are different to the ones for agents without this adaption.

Bibliography

- Ackerman, S. 1963. "Even-flow a scheduling method for reducing lateness in job shops." *Management Technology* 3: 20–32.
- Aydin, M. E., and E. Öztemel. 2000. "Dynamic job-shop scheduling using reinforcement learning agents." *Robotics and Autonomous Systems* 33 (2-3): 169–178.
- Bellemare, M. G., J. Veness, and M. Bowling. 2012. "Investigating contingency awareness using Atari 2600 games." In *Twenty-Sixth AAAI Conference on Artificial Intelligence*, .
- Benjamini, Y., and Y. Hochberg. 1995. "Controlling the false discovery rate: a practical and powerful approach to multiple testing." *Journal of the Royal statistical society: series B (Methodological)* 57 (1): 289–300.
- Bertrand, J. W. M. 1983. "The effect of workload dependent due-dates on job shop performance." *Management Science* 29 (7): 799–816.
- Bertsekas, D. P., D. P. Bertsekas, D. P. Bertsekas, and D. P. Bertsekas. 1995. *Dynamic programming and optimal control*. Vol. 1. Athena scientific Belmont, MA.
- Blackwell, D. 1962. "Discrete dynamic programming." *The Annals of Mathematical Statistics* 344: 719–726.
- Chang, P. C., Y. W. Wang, and C. J. Ting. 2008. "A fuzzy neural network for the flow time estimation in a semiconductor manufacturing factory." *International Journal of Production Research* 46 (4): 1017–1029.
- Chung, S.-H., and H.-W. Huang. 2002. "Cycle time estimation for wafer fab with engineering lots." *IIE Transactions* 34 (2): 105–118. <https://doi.org/10.1080/07408170208928854>.
- Conway, R., and W. Maxwell. 1967. *Theory of Scheduling*. Addison-Wesley Publ.
- Das, T. K., A. Gosavi, S. Mahadevan, and N. Marchallick. 1999. "Solving semi-Markov decision problems using average reward reinforcement learning." *Management Science* 45 (4): 560–574.
- Doltsinis, S., P. Ferreira, and N. Lohse. 2012. "Reinforcement learning for production ramp-up: A Q-batch learning approach." In *2012 11th International Conference on Machine Learning and Applications*, Vol. 1, 610–615. IEEE.
- Eilon, S., and I. G. Chowdhury. 1976. "Due dates in job shop scheduling." *The International Journal of Production Research* 14 (2): 223–237.
- Enns, S. T. 1995. "An integrated system for controlling shop loading and work flow." *International Journal of Production Research* 33 (10): 2801–2820.
- Enns, S. T., and P. Suwanruji. 2004. "Work load responsive adjustment of planned lead times." *Journal of Manufacturing Technology Management* 15 (1): 90–100.
- Gelders, L. F., and L. N. Van Wassenhove. 1982. "Hierarchical Integration in Production Planning: Theory and Practice." *Journal of Operations Management* 3 (1): 27–35.

- Gijsbrechts, J., R. N. Boute, J. A. Van Mieghem, and D. Zhang. 2018. “Can deep reinforcement learning improve inventory management? performance and implementation of dual sourcing-mode problems.” *Performance and Implementation of Dual Sourcing-Mode Problems (December 17, 2018)* .
- Glorot, X., and Y. Bengio. 2010. “Understanding the difficulty of training deep feedforward neural networks.” In *Proceedings of the thirteenth international conference on artificial intelligence and statistics*, 249–256.
- Howard, R. A. 1964. “Dynamic programming and Markov processes.” .
- Hoyt, J. 1978. “Dynamic lead times that fit today’s dynamic planning (QUOAT lead times).” *Production and Inventory Management* 19 (1): 63–71.
- Hsu, S. Y., and D. Y. Sha. 2004. “Due date assignment using artificial neural networks under different shop floor control strategies.” *International Journal of Production Research* 42 (9): 1727–1745.
- Kanet, J. J. 1986. “Toward a better understanding of lead times in MRP systems.” *Journal of Operations Management* 6 (3): 305–315.
- Kaplan, A. C., and A. T. Unal. 1993. “A probabilistic cost-based due date assignment model for job shops.” *International Journal of Production Research* 31 (12): 2817–2834.
- Kim, S.-C., and P. M. Bobrowski. 1995a. “Evaluating order release mechanisms in a job shop with sequence-dependent setup times.” *Production and Operations Management* 4 (2): 163–180.
- Kim, S.-C., and P. M. Bobrowski. 1995b. “Evaluating Order Release Mechanisms in a Job Shop With Sequence-Dependent Setup Times.” *Production and Operations Management* 4 (2): 163–180. <https://onlinelibrary.wiley.com/doi/abs/10.1111/j.1937-5956.1995.tb00048.x>.
- Kingma, D. P., and J. Ba. 2014. “Adam: A method for stochastic optimization.” *arXiv preprint arXiv:1412.6980* .
- Knollmann, M., and K. Windt. 2013. “Control-theoretic analysis of the Lead Time Syndrome and its impact on the logistic target achievement.” *Procedia CIRP* 7: 97–102.
- Law, A. M., and W. D. Kelton. 2000. *Simulation Modeling & Analysis*. 3rd ed. New York: McGraw-Hill, Inc.
- Lee, C.-Y., S. Piramuthu, and Y.-K. Tsai. 1997. “Job shop scheduling with a genetic algorithm and machine learning.” *International Journal of Production Research* 35 (4): 1171–1191. <https://doi.org/10.1080/002075497195605>.
- Loshchilov, I., and F. Hutter. 2017. “Decoupled weight decay regularization.” *arXiv preprint arXiv:1711.05101* .
- Mahadevan, S. 1996a. “An average-reward reinforcement learning algorithm for computing bias-optimal policies.” In *AAAI/IAAI, Vol. 1*, 875–880.
- Mahadevan, S. 1996b. “Average reward reinforcement learning: Foundations, algorithms, and empirical results.” *Machine Learning* 22: 159–195.
- Mahadevan, S. 1996c. “Optimality criteria in reinforcement learning.” In *Proceedings of the AAAI Fall Symposium on Learning Complex Behaviors in Adaptive Intelligent Systems*, .
- Mahadevan, S. 1996d. “Sensitive discount optimality: Unifying discounted and average reward reinforcement learning.” In *ICML*, 328–336.

- Mahadevan, S., N. Marchallick, T. K. Das, and A. Gosavi. 1997. "Self-improving factory simulation using continuous-time average-reward reinforcement learning." In *MACHINE LEARNING-INTERNATIONAL WORKSHOP THEN CONFERENCE*-, 202–210.
- Mahadevan, S., and G. Theocharous. 1998. "Optimizing Production Manufacturing Using Reinforcement Learning." In *FLAIRS Conference*, 372–377.
- Mather, H., and G. W. Plossl. 1978. "Priority fixation versus throughput planning." *Production and Inventory Management* 19: 27–51.
- Miller, B. L., and A. F. Veinott. 1969. "Discrete Dynamic Programming with a Small Interest Rate." *The Annals of Mathematical Statistics* 40 (2): 366–370. <http://www.jstor.org/stable/2239451>.
- Mnih, V., A. P. Badia, M. Mirza, A. Graves, T. Lillicrap, T. Harley, D. Silver, and K. Kavukcuoglu. 2016. "Asynchronous methods for deep reinforcement learning." In *International conference on machine learning*, 1928–1937.
- Mnih, V., K. Kavukcuoglu, D. Silver, A. A. Rusu, J. Veness, M. G. Bellemare, A. Graves, et al. 2015. "Human-level control through deep reinforcement learning." *Nature* 518 (7540): 529.
- Pahl, J., S. Voss, and D. L. Woodruff. 2007. "Production planning with load dependent lead times: an update of research." *Annals of Operations Research* 153 (1): 297–345. <GotoISI>://WOS:000247203300013.
- Park, P. S., and P. M. Bobrowski. 1989. "Job release and labor flexibility in a dual resource constrained job shop." *Journal of Operations Management* 8 (3): 230–249.
- Paternina-Arboleda, C. D., and T. K. Das. 2001. "Intelligent dynamic control policies for serial production lines." *Iie Transactions* 33 (1): 65–77.
- Patil, R. 2008. "Using ensemble and metaheuristics learning principles with artificial neural networks to improve due date prediction performance." *International Journal of Production Research* 46 (21): 6009–6027.
- Philipoom, P. R., L. P. Rees, and L. Wiegmann. 1994. "Using Neural Networks to Determine Internally-Set Due-Date Assignments for Shop Scheduling*." *Decision Sciences* 25 (5-6): 825–851.
- Philipoom, P. R., L. Wiegmann, and L. P. Rees. 1997. "Cost-based due-date assignment with the use of classical and neural-network approaches." *Naval Research Logistics (NRL)* 44 (1): 21–46.
- Puterman, M. L. 1994. "Markov Decision Processes. J." *Wiley and Sons* .
- Ragatz, G. J., and V. A. Mabert. 1988. "An evaluation of order release mechanisms in a job-shop environment." *Decision Sciences* 19: 167–189.
- Schneckenreither, M. 2019. "Blackwell-Optimal Reinforcement Learning." Working paper.
- Schneckenreither, M. 2020. "Average Reward Adjusted Discounted Reinforcement Learning: Near-Blackwell-Optimal Policies for Real-World Applications." *arXiv preprint arXiv:2004.00857* .
- Schneckenreither, M., and S. Haeussler. 2019. *Reinforcement Learning Methods for Operations Research Applications: The Order Release Problem*, 46. Springer International Publishing.

- Schneckenreither, M., S. Haeussler, and C. Gerhold. 2020. "Order release planning with predictive lead times: a machine learning approach." *International Journal of Production Research* .
- Schneeweiß, C. 1995. "Hierarchical structures in organisations: A conceptual framework." *European Journal of Operational Research* 86 (1): 4–31.
- Schwartz, A. 1993. "A reinforcement learning method for maximizing undiscounted rewards." In *Proceedings of the tenth international conference on machine learning*, Vol. 298, 298–305.
- Selcuk, B. 2013. "Adaptive lead time quotation in a pull production system with lead time responsive demand." *Journal of Manufacturing Systems* 32 (1): 138–146.
- Selcuk, B., I. J. B. F. Adan, T. G. de Kok, and J. C. Fransoo. 2009. "An explicit analysis of the lead time syndrome: stability condition and performance evaluation." *International Journal of Production Research* 47 (9): 2507–2529.
- Selcuk, B., J. C. Fransoo, and T. G. De Kok. 2006. "The effect of updating lead times on the performance of hierarchical planning systems." *International Journal of Production Economics* 104 (2): 427–440.
- Silver, D., A. Huang, C. J. Maddison, A. Guez, L. Sifre, G. Van Den Driessche, J. Schrittwieser, et al. 2016. "Mastering the game of Go with deep neural networks and tree search." *nature* 529 (7587): 484.
- Silver, D., T. Hubert, J. Schrittwieser, I. Antonoglou, M. Lai, A. Guez, M. Lanctot, et al. 2017. "Mastering chess and shogi by self-play with a general reinforcement learning algorithm." *arXiv preprint arXiv:1712.01815* .
- Sutton, R. S., A. G. Barto, et al. 1998. *Introduction to reinforcement learning*. Vol. 2. MIT press Cambridge.
- Tadepalli, P., and D. Ok. 1998. "Model-based average reward reinforcement learning." *Artificial intelligence* 100 (1-2): 177–224.
- Tai, Y. T., W. L. Pearn, and J. H. Lee. 2012. "Cycle time estimation for semiconductor final testing processes with Weibull-distributed waiting time." *International Journal of Production Research* 50 (2): 581–592.
- Tatsiopoulos, I. P., and B. G. Kingsman. 1983. "Lead time management." *European Journal of Operational Research* 14 (4): 351–358.
- Veinott, A. F. 1969. "Discrete dynamic programming with sensitive discount optimality criteria." *The Annals of Mathematical Statistics* 40 (5): 1635–1660.
- Vig, M. M., and K. J. Dooley. 1991. "Dynamic rules for due-date assignment." *The International Journal of Production Research* 29 (7): 1361–1377.
- Wang, Y.-C., and J. M. Usher. 2005. "Application of reinforcement learning for agent-based production scheduling." *Engineering Applications of Artificial Intelligence* 18 (1): 73–82.
- Wang, Y.-C., and J. M. Usher. 2007. "A reinforcement learning approach for developing routing policies in multi-agent production scheduling." *The International Journal of Advanced Manufacturing Technology* 33 (3-4): 323–333.
- Waschneck, B., A. Reichstaller, L. Belzner, T. Altenmüller, T. Bauernhansl, A. Knapp, and A. Kyek. 2018. "Optimization of global production scheduling with deep reinforcement learning." *Procedia CIRP* 72 (1): 1264–1269.

- Wiendahl, H. P. 1995. *Load-Oriented Manufacturing Control*. 1st ed. Berlin: Springer.
- Wisner, J. D. 1995. “A review of the order release policy research.” *International Journal of Operations & Production Management* .
- Zäpfel, G. 1982. *Produktionswirtschaft: Operatives Produktions-Management*. de Gruyter.
- Zhang, W., and T. G. Dietterich. 1995. “A reinforcement learning approach to job-shop scheduling.” In *IJCAI*, Vol. 95, 1114–1120. Citeseer.

A Deep Q-Learning Algorithm

Algorithm 3 presents the detailed deep Q-Learning algorithm. It is very similar to the established ARAL algorithm that is given in Algorithm 2. However, it does not compute the average reward. Furthermore, it also operates on a single estimation of the overall future discounted reward.

Algorithm 3 Deep Q-Learning

- 1: Initialise state s_0 and network parameters θ_X^T, θ_X^W randomly, set exploration rate $0 < p_{\text{expl}} \leq 1$, target network adaption rate $0 < \gamma \ll 1$ and discount factor $0 < \gamma_1 < 1$
 - 2: **while** the stopping criterion is not fulfilled **do**
 - 3: **for each** independent agent $i \in \text{ag}_1, \text{ag}_2, \dots, \text{ag}_n$ **do**
 - 4: With probability $\frac{p_{\text{expl}}}{n}$ choose action $a_{t,i}$ randomly or otherwise the one fulfilling $\min_{a \in \mathcal{A}(i)} X_{\gamma_1}^\pi(s_t, a; \theta_X^W)$.
 - 5: Carry out action $a_t = (a_{t,\text{ag}_1}, \dots, a_{t,\text{ag}_n})$, observe reward r_t and resulting state s_{t+1} . Store experience $(s_t, a_t, a_t^{\text{rd}}, r_t, s_{t+1})$ in the experience replay memory M .
 - 6: **if** $t \bmod n_{\text{step}} \equiv 0$ **then**
 - 7: Reset gradient: $d\theta_X^W \leftarrow 0$
 - 8: Sample m random mini-batches of n_{step} consecutive experiences from M numbered by $1 \leq i \leq n_{\text{step}}$, each with $(s_i, a_i^{\text{rd}}, a_i, r_i, s_{i+1})$.
 - 9: **for each** mini-batch **do**
 - 10: **for each** experience, starting with the latest seen experience n_{step} **do**
 - 11: **if** experience n_{step} or $\neg a_i^{\text{rd}}$ or $p_{\text{expl}} > p_{\text{learn}}$ **then**
 - 12: $\text{target}_{\gamma_1} \leftarrow \min_a X_{\gamma_1}^\pi(s_{i+1}, a; \theta_X^T)$
 - 13: $\text{target}_{\gamma_1} \leftarrow r_i + \gamma_1 \text{target}_{\gamma_1}$
 - 14: $d\theta_X^W \leftarrow d\theta_X^W + \partial(\text{target}_{\gamma_1} - X_{\gamma_1}^\pi(s_{i+1}, a; \theta_X^W))^2 / \partial \theta_X^W$
 - 15: Perform update of θ_X^W using $d\theta_X^W / (m \cdot n_{\text{step}})$
 - 16: Every C steps exponential smoothly set target network: $\theta_X^T \leftarrow (1 - \gamma)\theta_X^T + \gamma\theta_X^W$
 - 17: Set $s \leftarrow s'$, $t \leftarrow t + 1$ and decay parameters
-

B Statistical Analysis of the Results

Table 4 and Table 5 provide the statistical results of the computation results. We used the Friedman test with a significance level of $p = 0.05$ for a statistical analysis of the mean sum of costs. As expected the omnibus null hypotheses (all samples are from the same distribution) are rejected for all cases. Therefore, we conducted pairwise Conover post-hoc tests adjusted by the Benjamini-Hochberg FDR method (Benjamini and Hochberg, 1995) to reduce liberality. These results are shown in Table 4 and Table 5 for the uniformly and exponentially distributed interarrival times of the demand respectively. The values given in bold are below the significance level of 5% and therefore indicate which row-column tuples are not statistically significant.

70- \mathcal{U}	QL _{0.99}	ARAL _{0.99}	ARAL _{1.00}	BIL1	BIL2	BIL3
ARAL _{0.99}	3.37e-04					
ARAL _{1.00}	3.55e-01	1.04e-05				
BIL1	1.94e-45	2.48e-51	6.90e-44			
BIL2	1.09e-24	3.35e-16	9.57e-27	3.66e-64		
BIL3	1.78e-20	9.27e-29	2.44e-18	1.35e-22	6.26e-47	
BIL4	6.30e-34	3.85e-41	4.83e-32	1.16e-08	3.00e-56	1.16e-08
80- \mathcal{U}	QL _{0.99}	ARAL _{0.99}	ARAL _{1.00}	BIL1	BIL2	BIL3
ARAL _{0.99}	4.78e-16					
ARAL _{1.00}	5.26e-04	3.96e-08				
BIL1	1.02e-40	1.98e-21	3.92e-34			
BIL2	6.98e-10	7.12e-03	2.06e-03	1.31e-27		
BIL3	2.16e-02	8.93e-11	2.03e-01	1.57e-36	2.46e-05	
BIL4	6.16e-33	1.10e-11	2.74e-25	5.33e-05	5.38e-18	5.30e-28
90- \mathcal{U}	QL _{0.99}	ARAL _{0.99}	ARAL _{1.00}	BIL1	BIL2	BIL3
ARAL _{0.99}	5.18e-01					
ARAL _{1.00}	1.94e-04	1.56e-05				
BIL1	8.40e-42	9.61e-43	1.10e-34			
BIL2	2.36e-33	1.10e-34	5.21e-25	1.56e-05		
BIL3	4.16e-04	3.64e-05	8.19e-01	6.27e-35	1.81e-25	
BIL4	2.33e-20	6.57e-22	1.79e-11	2.93e-18	4.39e-08	5.97e-12

Table 4: P-Values from the statistical analysis for the results of moderate variability in the system. Tuples marked bold are not significantly different.

70-Exp	QL _{0.99}	ARAL _{0.99}	ARAL _{1.00}	BIL1	BIL2	BIL3
ARAL _{0.99}	1.75e-02					
ARAL _{1.00}	6.15e-07	5.50e-12				
BIL1	7.62e-35	1.74e-39	1.32e-23			
BIL2	3.33e-01	1.02e-03	3.17e-05	6.06e-33		
BIL3	5.96e-17	1.79e-11	6.09e-29	7.00e-52	3.47e-19	
BIL4	1.16e-24	6.25e-30	1.68e-12	4.68e-06	1.56e-22	3.43e-44
80-Exp	QL _{0.99}	ARAL _{0.99}	ARAL _{1.00}	BIL1	BIL2	BIL3
ARAL _{0.99}	7.07e-02					
ARAL _{1.00}	1.74e-03	1.72e-01				
BIL1	2.39e-35	6.77e-39	2.38e-41			
BIL2	1.22e-25	1.00e-29	1.21e-32	1.13e-05		
BIL3	2.72e-05	1.33e-02	2.44e-01	3.65e-43	5.68e-35	
BIL4	2.07e-13	1.12e-17	6.55e-21	3.62e-17	6.09e-07	1.45e-23
90-Exp	QL _{0.99}	ARAL _{0.99}	ARAL _{1.00}	BIL1	BIL2	BIL3
ARAL _{0.99}	2.61e-06					
ARAL _{1.00}	5.03e-13	1.74e-03				
BIL1	9.64e-23	2.73e-33	2.22e-39			
BIL2	3.46e-12	1.32e-23	1.30e-30	1.23e-05		
BIL3	6.86e-01	1.23e-05	3.46e-12	1.32e-23	5.03e-13	
BIL4	2.45e-04	2.38e-01	2.60e-05	6.39e-31	6.30e-21	9.46e-4

Table 5: P-Values from the statistical analysis for the results of high variability in the system. Tuples marked bold are not significantly different.